

Spatial: A Language and Compiler for Application Accelerators

David Koeplinger

Matthew
Feldman

Raghu
Prabhakar

Yaqi
Zhang

Stefan
Hadjis

Tian
Zhao

Ruben
Fiszel

Luigi
Nardi

Ardavan
Pedram

Christos
Kozyrakis

Kunle
Olukotun

Spatial Resources

- Language documentation and tutorials:
 - spatial.stanford.edu
- Spatial Google Group:
 - groups.google.com/forum/#!forum/spatial-lang-users
- Spatial Github Repo:
 - github.com/stanford-ppl/spatial-lang

Increasing Demand for Reconfigurability

There is growing demand for **reconfigurable architectures** as **application accelerators** in data centers for performance and energy efficiency

Improvements in Performance, Energy Efficiency



Software Defined Accelerators for DNNs

Distributed real-time deep learning on FPGAs
[HotChips '14]

2 – 4x perf vs. CPU/GPU
2 – 3x perf/W vs. CPU/GPU

XPU: Programmable FPGA Accelerator

[HotChips '17]

64x perf vs. CPU
25x perf/W vs. CPU



Catapult “Configurable Cloud”

Distributed application and network acceleration

“...already been deployed at **hyperscale** and is how most new Microsoft data center servers are configured.” [Caulfield et. al. Micro '16]

~10x perf vs. CPU
~ 9x perf/W vs. CPU

Project Brainwave

Distributed real-time deep learning on FPGAs

[HotChips '17]




CGRA for Parallel Patterns


[ISCA '17]

Up to 95x perf vs. FPGA
Up to 75x perf/W vs. FPGA

Accessing Reconfigurable Architectures



Accessing Reconfigurable Architectures



Key Programming Challenges

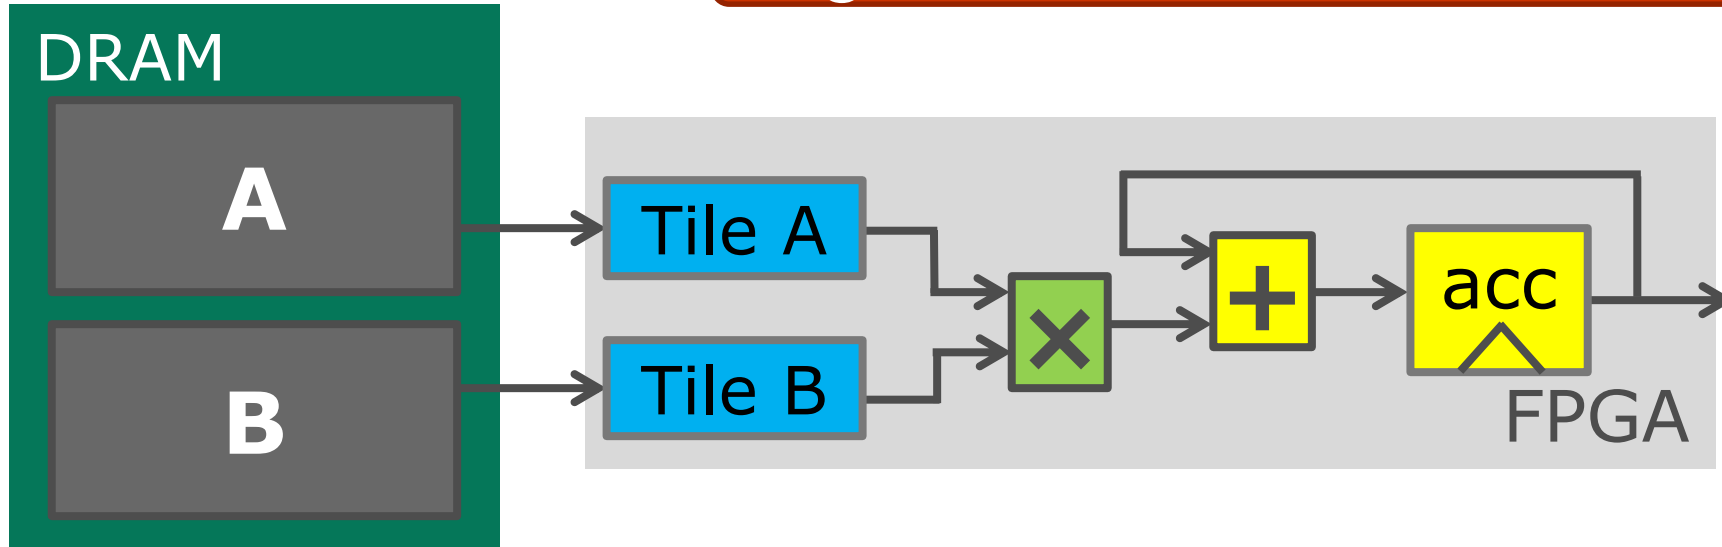
- Pipeline timing
 - Fine-grained timing of control signals
 - Management of multiple clock domains
- Disjoint memory spaces
 - No hardware-managed cache
 - Requires manual data partitioning and timing of memory operations
- Limited compute and memory resources
- Huge parameter design spaces
 - Grows exponentially - even relatively small designs can have very large spaces
 - Parameters are interdependent and can change runtime by orders of magnitude
 - Manual exploration is tedious and usually suboptimal

Key Programming Challenges (Continued..)


- Quick verification
 - Hardware synthesis is too slow to be part of the debug loop
 - Writing test benches is challenging and requires waveform debugging
- Managing host-accelerator interface
 - Tile transfers, register interfaces, data streams, etc.
- Incorporating IP cores into design

Design Space Example

Algorithm: Dot Product of Vectors A and B

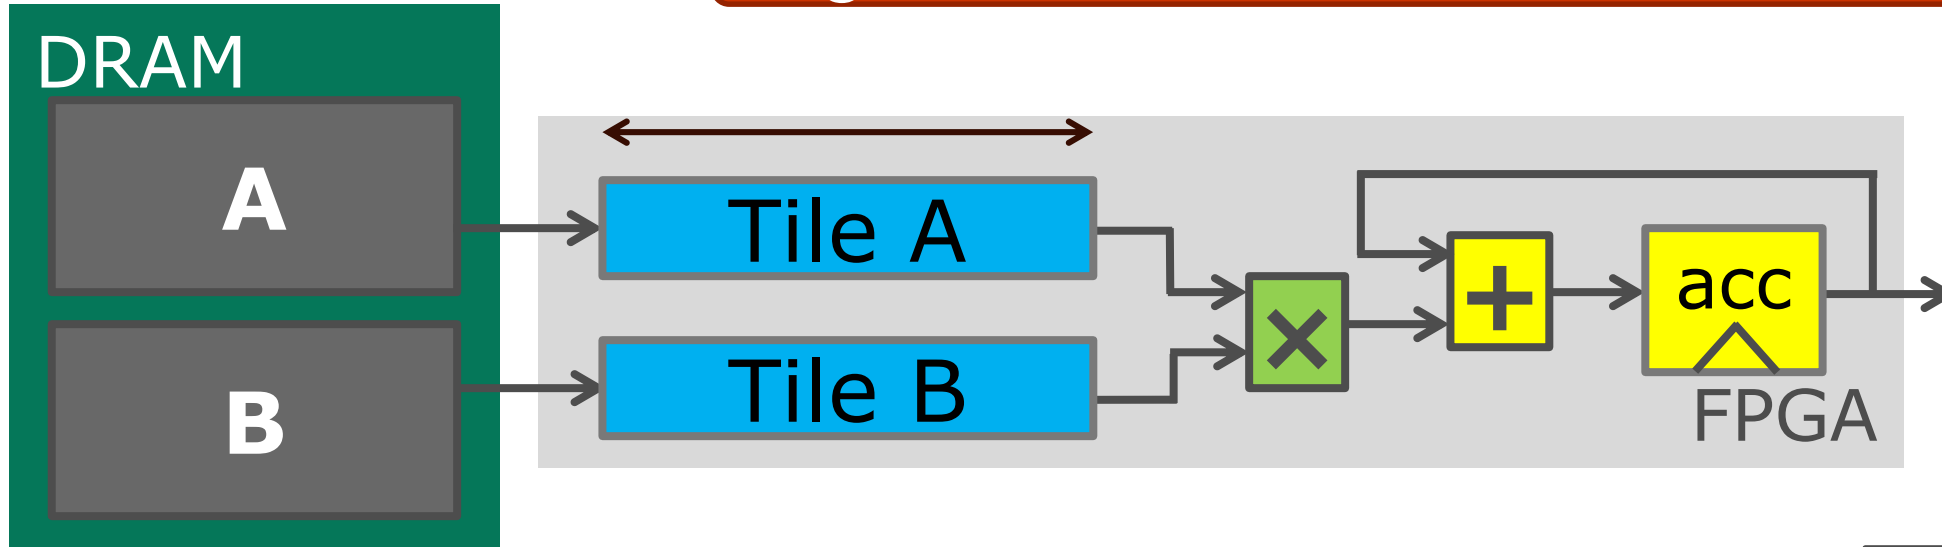


Small and simple, but slow!




Important Parameters: Tile Sizes

Algorithm: Dot Product of Vectors A and B



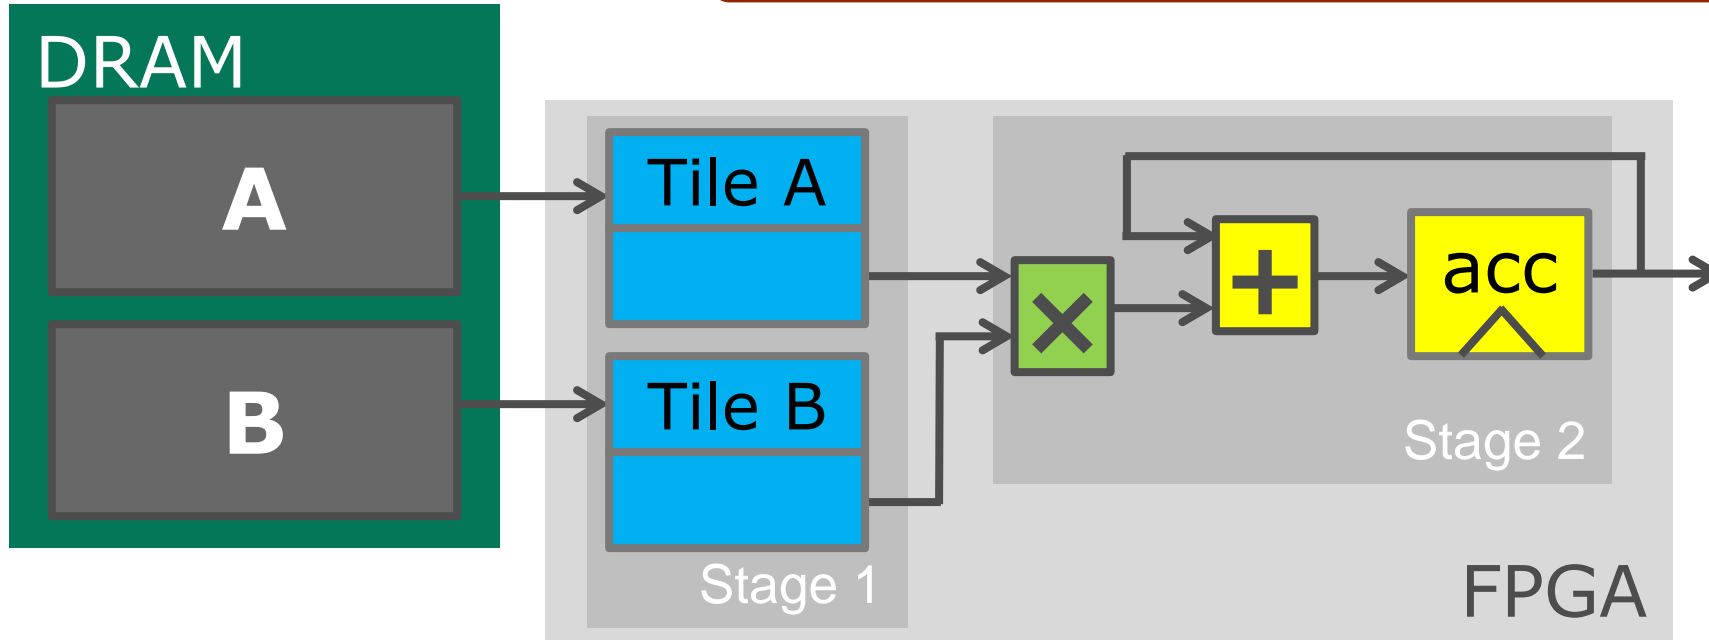
- Increases length of DRAM accesses
- Increases exploited spatial locality
- Increases local memory sizes




↓ Runtime
↓ Runtime
↑ Area




Important Parameters: Pipelining

Algorithm: Dot Product of Vectors A and B

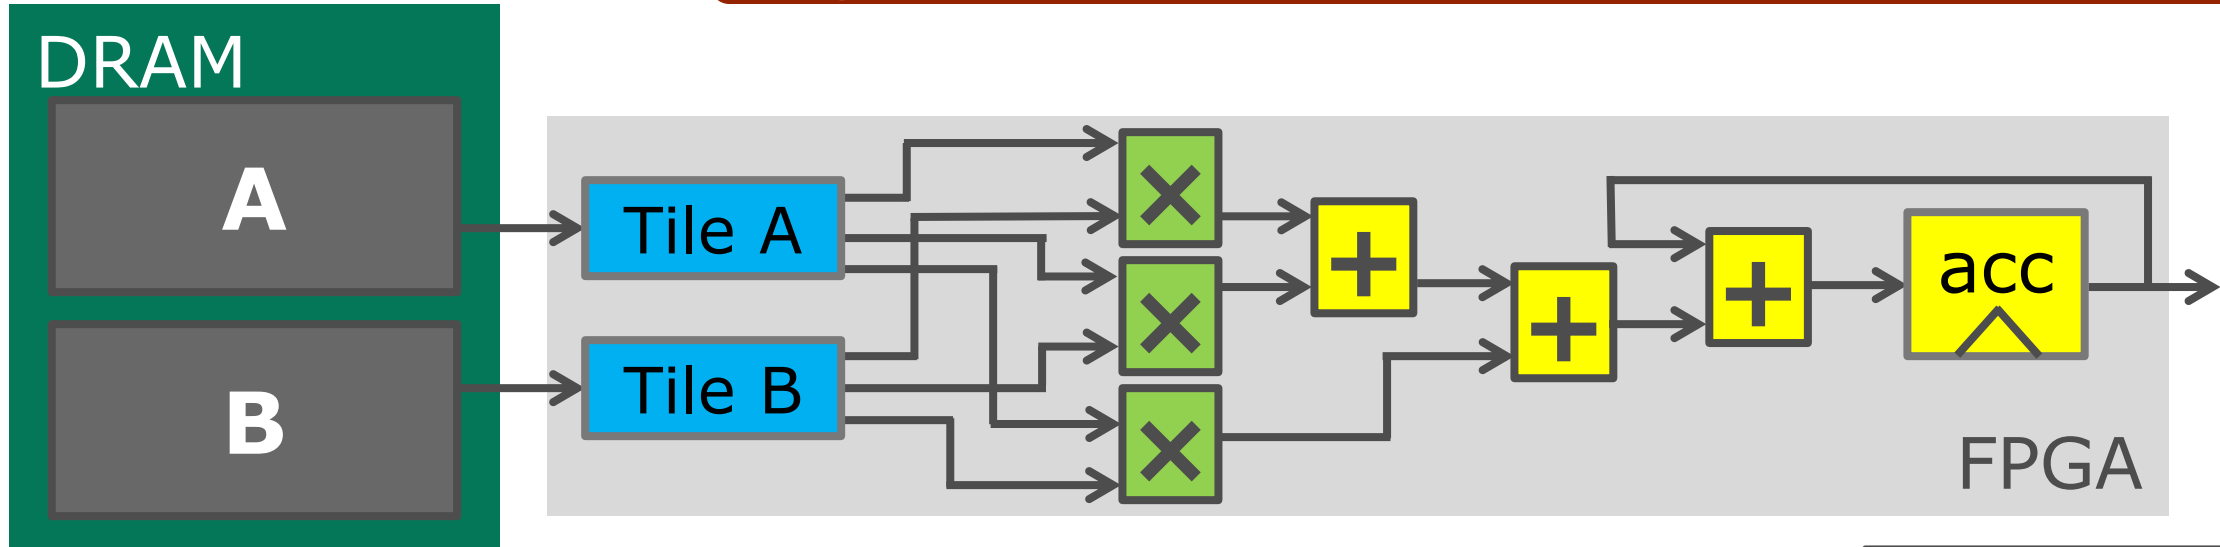


- Overlaps memory and compute  Runtime
- Increases local memory sizes  Area
- Adds synchronization logic  Area




Important Parameters: Parallelization

Algorithm: Dot Product of Vectors A and B



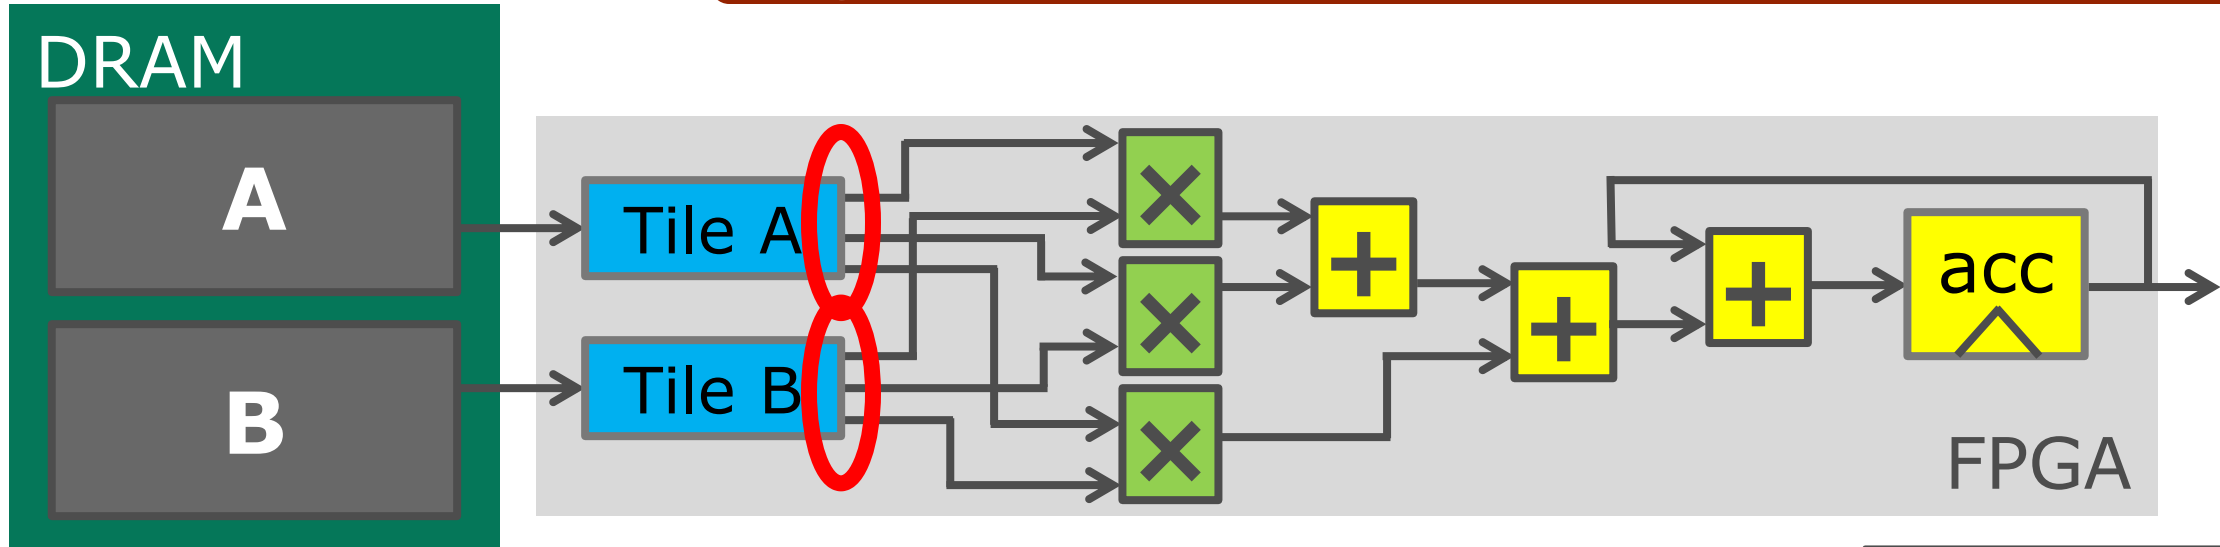
- Improves element throughput
- Duplicates compute resources

↓ Runtime
↑ Area




Important Parameters: Memory Banking

Algorithm: Dot Product of Vectors A and B




- Improves element throughput
- May duplicate memory resources

↓ Runtime
↑ Area



Language Requirements

1. Performs hardware optimizations for higher level frameworks
2. Productive language for “power” users (hardware programmers)
3. Produces efficient hardware



Language Comparisons

HDLs (Verilog, VHDL, Chisel, etc.)


Performance

✓ Arbitrary RTL

Productivity

✗ No high-level abstractions

Portability



High Level Synthesis (C, OpenCL)

Performance

✗ No memory hierarchy


✗ No arbitrary pipelining

Productivity

✓ Nested loops

✗ Difficult to tune

Portability



Spatial

Performance

✓ Memory hierarchy

✓ Arbitrary pipelining

✗ Can't write arbitrary RTL

Productivity

✓ Nested loops


✓ Automatic memory banking/buffering

✓ Automated design tuning

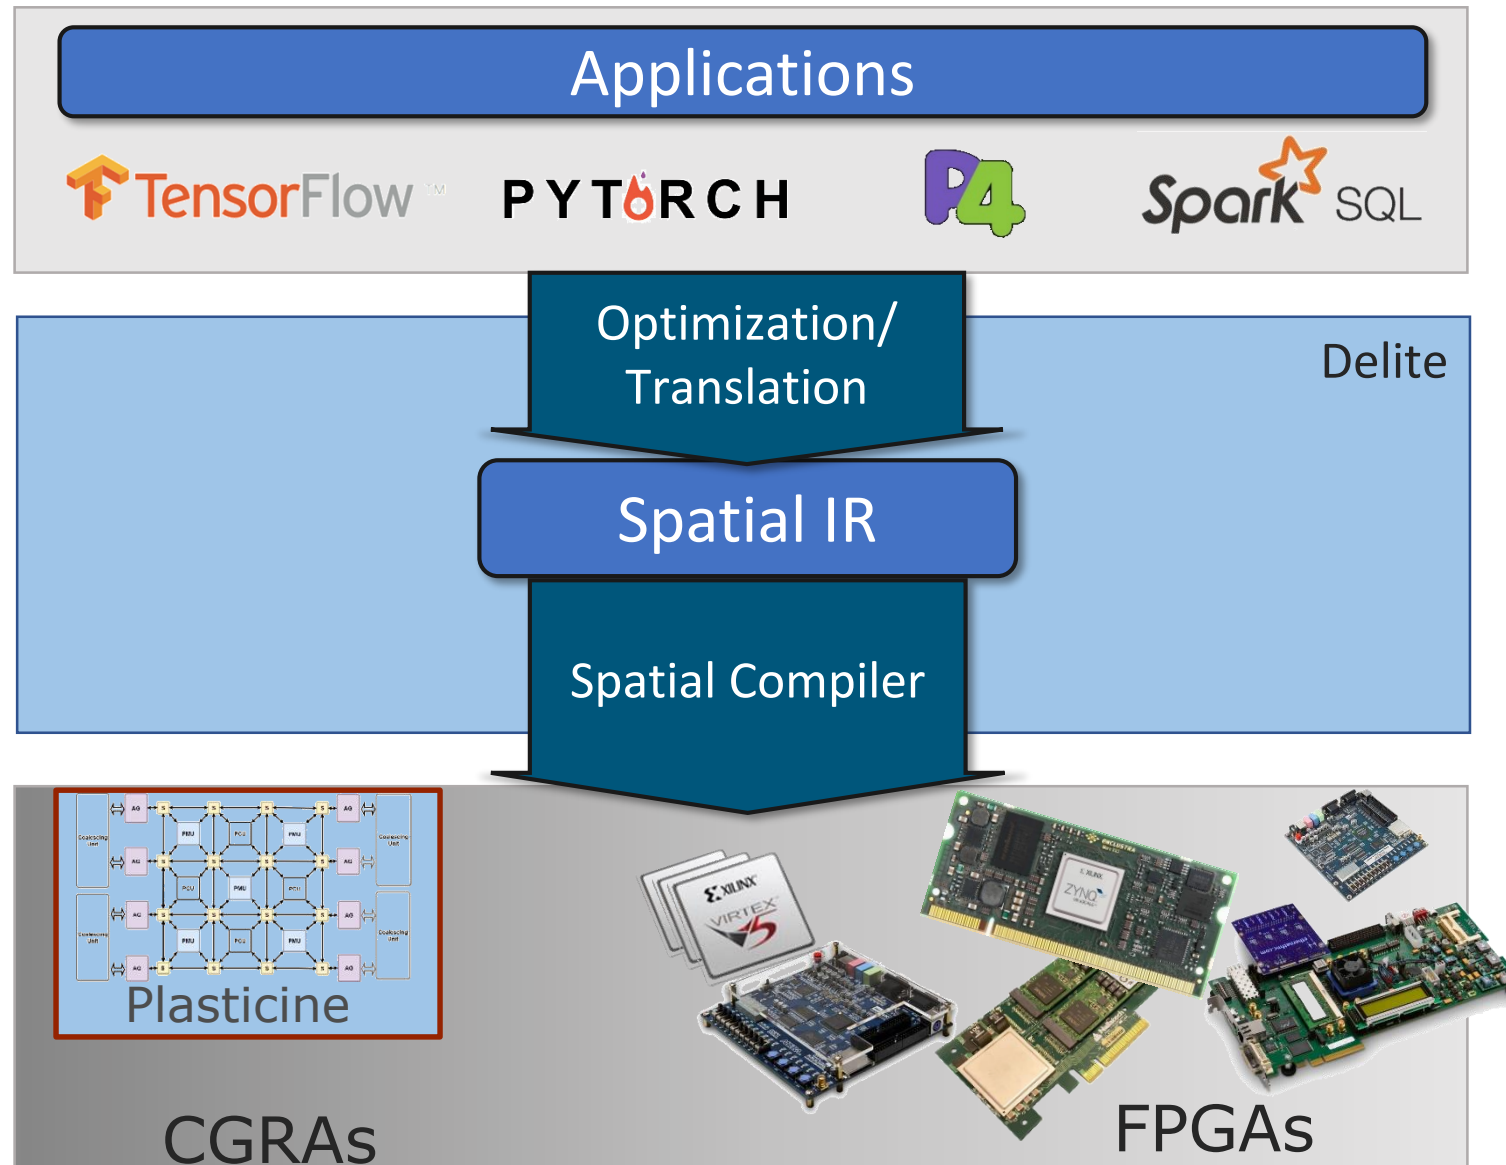
✓ Simple host communication API

Portability

✓ Target-generic source



Spatial as an Intermediate Representation



The Spatial Language: Memory Abstractions

Typed storage templates

```
val accum  = Reg[Double]  
val fifo   = FIFO[Float](D)  
val lbuf   = LineBuffer[Int](R,C)  
val pixels = ShiftReg[UInt8](R,C)
```

Explicit memory hierarchy

```
val buffer = SRAM[UInt8](C)  
val image  = DRAM[UInt8](H,W)
```

Explicit transfers across memory hierarchy
Dense and sparse access

```
buffer load image(i, j::j+C)  
out := x + y  
buffer gather image(a, 10)
```

Streaming abstractions

```
val videoIn  = StreamIn[RGB]  
val videoOut = StreamOut[RGB]
```

The Spatial Language: Control Abstractions

Blocking/non-blocking
interaction with host

```
Accel { ... }
```

```
Accel(*) { ... }
```

Arbitrary state machine / loop nesting
with implicit control signals

```
FSM[Int]{s => s != DONE }{  
  case STATE0 =>  
    Foreach(C by 1){j => ... }  
  case STATE1 => ...  
    Reduce(0)(C by 1){i => ... }  
}  
}{s => nextState(s) }
```

The Spatial Language: Design Parameters

Spatial templates capture a variety of design parameters:

Implicit/Explicit parallelization factors

```
val P = 16 (1 → 32)
Reduce(0)(N by 1 par P){i =>
  data(i)
}{(a,b) => a + b}
```

Implicit/Explicit control schemes

```
Stream.Foreach(0 until N){i =>
  ...
}
```

Explicit size parameters for stride
and buffer sizes

```
val B = 64 (64 → 1024)
val buffer = SRAM[Float](B)
Foreach(N by B){i =>
  ...
}
```

Implicit memory banking and buffering
schemes for parallelized access

```
Foreach(64 par 16){i =>
  buffer(i) // Parallel read
}
```


Dot Product in Spatial

```
val output  = ArgOut[Float]
val vectorA = DRAM[Float](N)
val vectorB = DRAM[Float](N)
```

```
Accel {
  Reduce(output)(N by B){ i =>
    val tileA = SRAM[Float](B)
    val tileB = SRAM[Float](B)
    val acc   = Reg[Float]

    tileA load vectorA(i :: i+B)
    tileB load vectorB(i :: i+B)

    Reduce(acc)(B by 1){ j =>
      tileA(j) * tileB(j)
    }{a, b => a + b}
  }{a, b => a + b}
}
```




Dot Product in Spatial

```
val output  = ArgOut[Float]
val vectorA = DRAM[Float](N)
val vectorB = DRAM[Float](N)
```

```
Accel {
  Reduce(output)(N by B){ i =>
    val tileA = SRAM[Float](B)
    val tileB = SRAM[Float](B)
    val acc    = Reg[Float]

    tileA load vectorA(i :: i+B)
    tileB load vectorB(i :: i+B)

    Reduce(acc)(B by 1){ j =>
      tileA(j) * tileB(j)
    }{a, b => a + b}
  }{a, b => a + b}
}
```




Dot Product in Spatial

```
val output  = ArgOut[Float]
val vectorA = DRAM[Float](N)
val vectorB = DRAM[Float](N)
```

```
Accel {
  Reduce(output)(N by B){ i =>
    val tileA = SRAM[Float](B)
    val tileB = SRAM[Float](B)
    val acc    = Reg[Float]

    tileA load vectorA(i :: i+B)
    tileB load vectorB(i :: i+B)

    Reduce(acc)(B by 1){ j =>
      tileA(j) * tileB(j)
    }{a, b => a + b}
  }{a, b => a + b}
}
```




Dot Product in Spatial

```
val output  = ArgOut[Float]
val vectorA = DRAM[Float](N)
val vectorB = DRAM[Float](N)
```

```
Accel {
  Reduce(output)(N by B){ i =>
    val tileA = SRAM[Float](B)
    val tileB = SRAM[Float](B)
    val acc    = Reg[Float]

    tileA load vectorA(i :: i+B)
    tileB load vectorB(i :: i+B)

    Reduce(acc)(B by 1){ j =>
      tileA(j) * tileB(j)
    }{a, b => a + b}
  }{a, b => a + b}
}
```



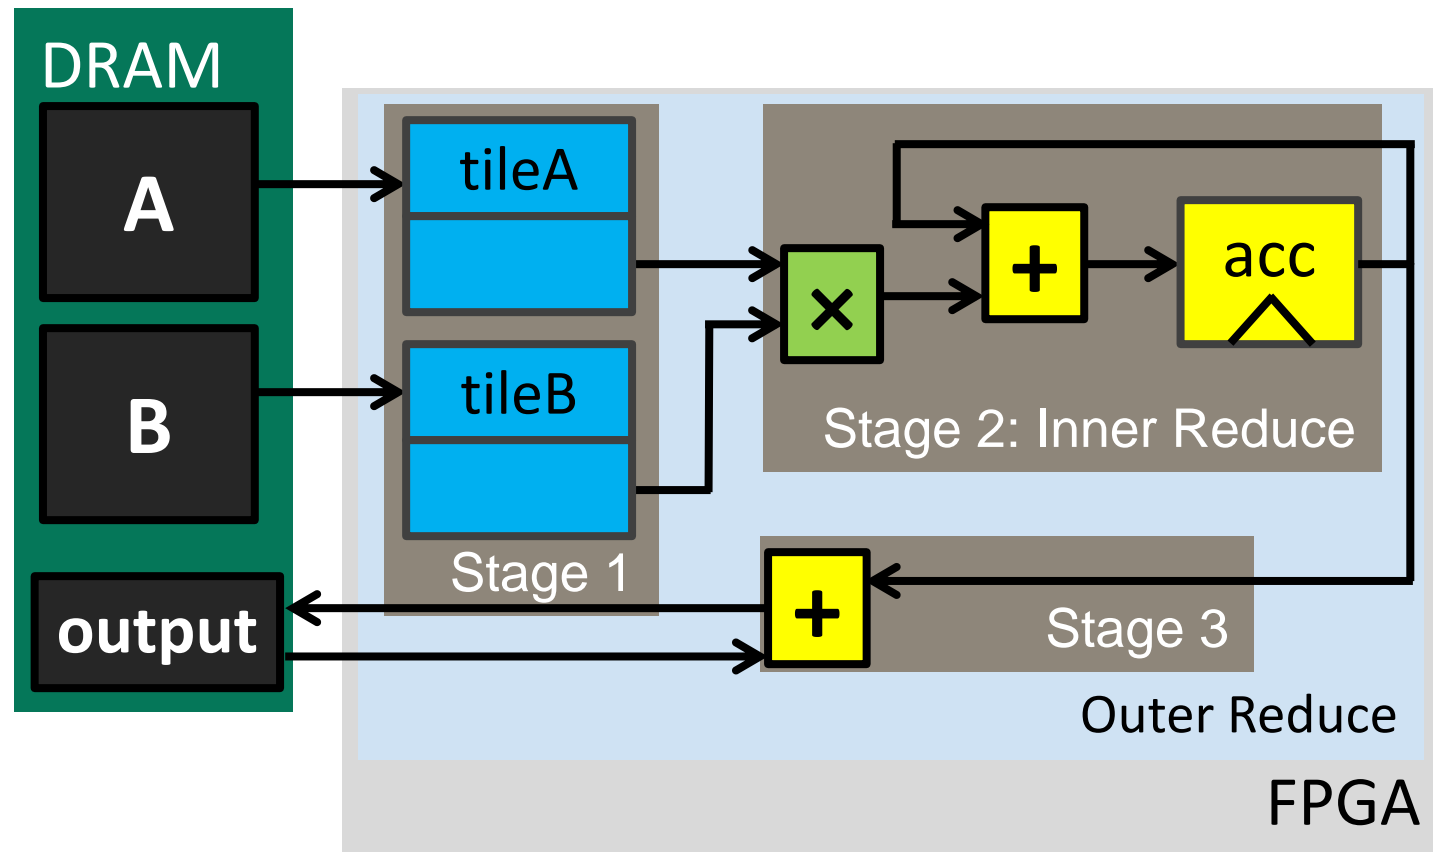
Dot Product in Spatial

```
val output  = ArgOut[Float]
val vectorA = DRAM[Float](N)
val vectorB = DRAM[Float](N)
```

```
Accel {
  Reduce(output)(N by B){ i =>
    val tileA = SRAM[Float](B)
    val tileB = SRAM[Float](B)
    val acc    = Reg[Float]

    tileA load vectorA(i :: i+B)
    tileB load vectorB(i :: i+B)

    Reduce(acc)(B by 1){ j =>
      tileA(j) * tileB(j)
    }{a, b => a + b}
  }{a, b => a + b}
}
```




Dot Product in Spatial

```
val output  = ArgOut[Float]
val vectorA = DRAM[Float](N)
val vectorB = DRAM[Float](N)
```

```
Accel {
  Reduce(output)(N by B){ i =>
    val tileA = SRAM[Float](B)
    val tileB = SRAM[Float](B)
    val acc   = Reg[Float]

    tileA load vectorA(i :: i+B)
    tileB load vectorB(i :: i+B)

    Reduce(acc)(B by 1){ j =>
      tileA(j) * tileB(j)
    }{a, b => a + b}
  }{a, b => a + b}
}
```




Dot Product in Spatial

```
val output  = ArgOut[Float]
val vectorA = DRAM[Float](N)
val vectorB = DRAM[Float](N)
```

```
Accel {
  Reduce(output)(N by B){ i =>
    val tileA = SRAM[Float](B)
    val tileB = SRAM[Float](B)
    val acc    = Reg[Float]

    tileA load vectorA(i :: i+B)
    tileB load vectorB(i :: i+B)

    Reduce(acc)(B by 1){ j =>
      tileA(j) * tileB(j)
    }{a, b => a + b}
  }{a, b => a + b}
}
```



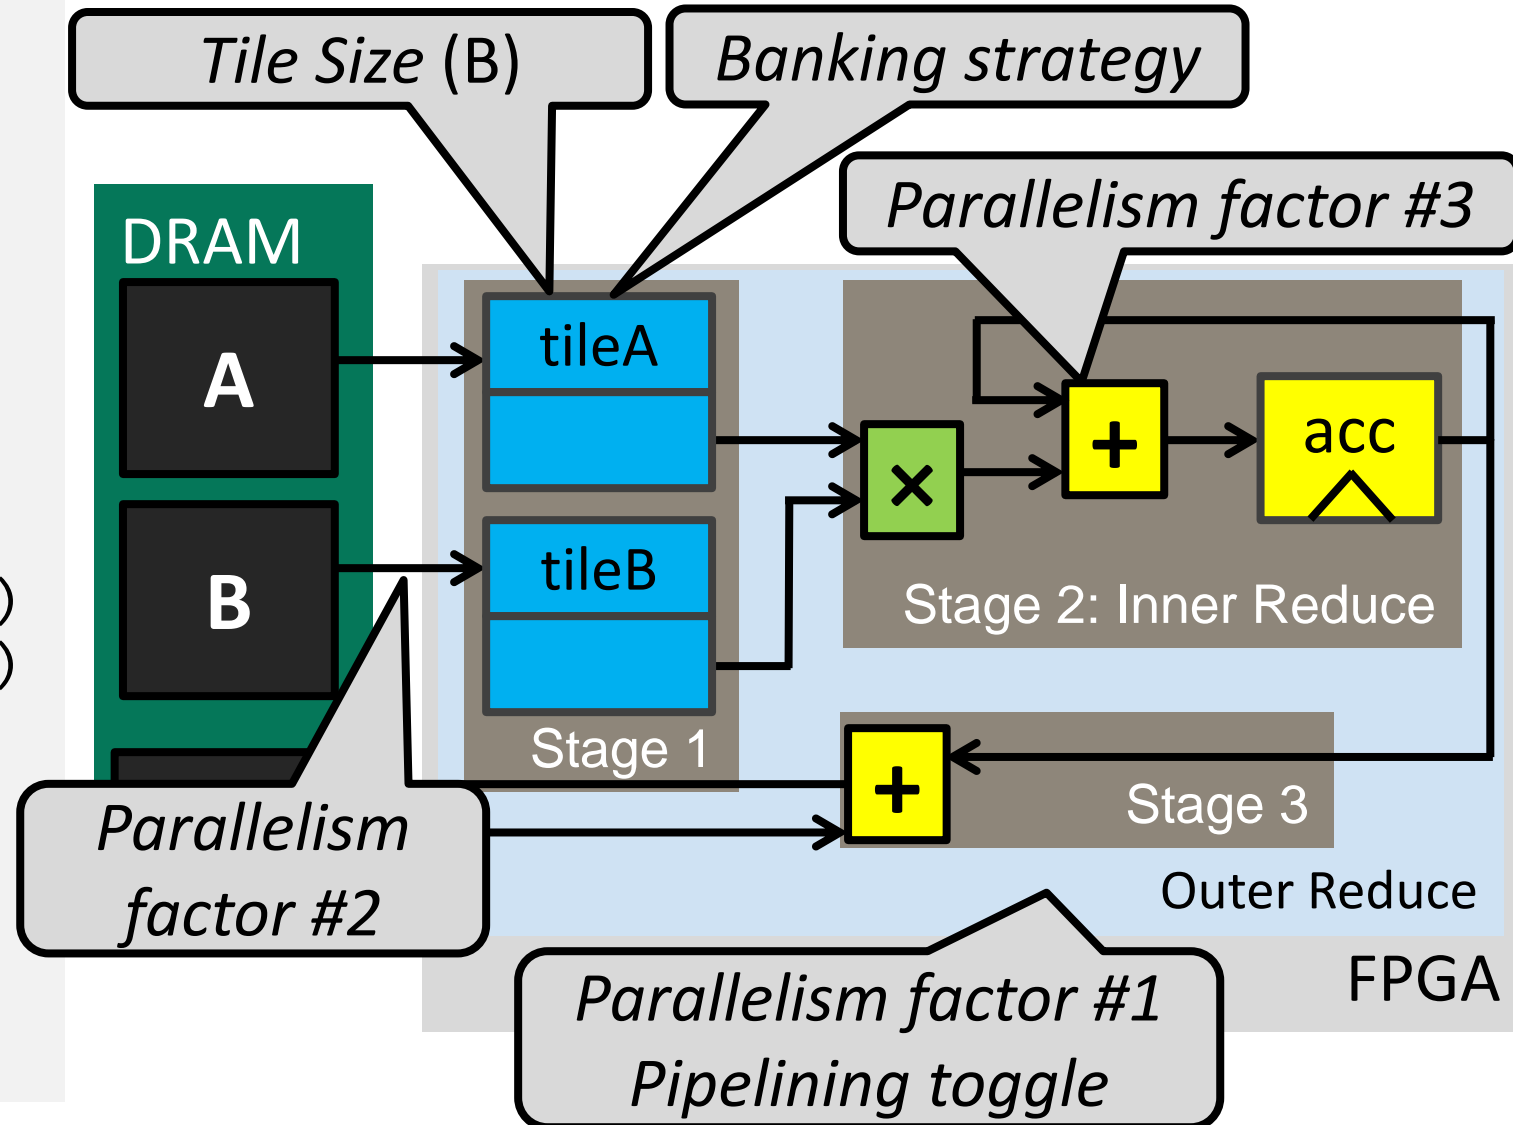
Dot Product in Spatial

```
val output  = ArgOut[Float]
val vectorA = DRAM[Float](N)
val vectorB = DRAM[Float](N)
```

```
Accel {
  Reduce(output)(N by B){ i =>
    val tileA = SRAM[Float](B)
    val tileB = SRAM[Float](B)
    val acc    = Reg[Float]

    tileA load vectorA(i :: i+B)
    tileB load vectorB(i :: i+B)

    Reduce(acc)(B by 1){ j =>
      tileA(j) * tileB(j)
    }{a, b => a + b}
  }{a, b => a + b}
}
```



Design Tuning Performance


Spatial:

Benchmark	Designs	Search Time
Dot Product	5,426	5.3 ms / design
Outer Product	1,702	30 ms / design
TPCHQ6	5,426	2.2 ms / design
Blackscholes	6500x Speedup Over HLS!	27 ms / design
Matrix Multiply		1.1 ms / design
K-Means		20 ms / design
GDA	42,800	17 ms / design

Vivado HLS:

	Designs	Search Time
GDA	250	1.85 min / design

GDA Design Space



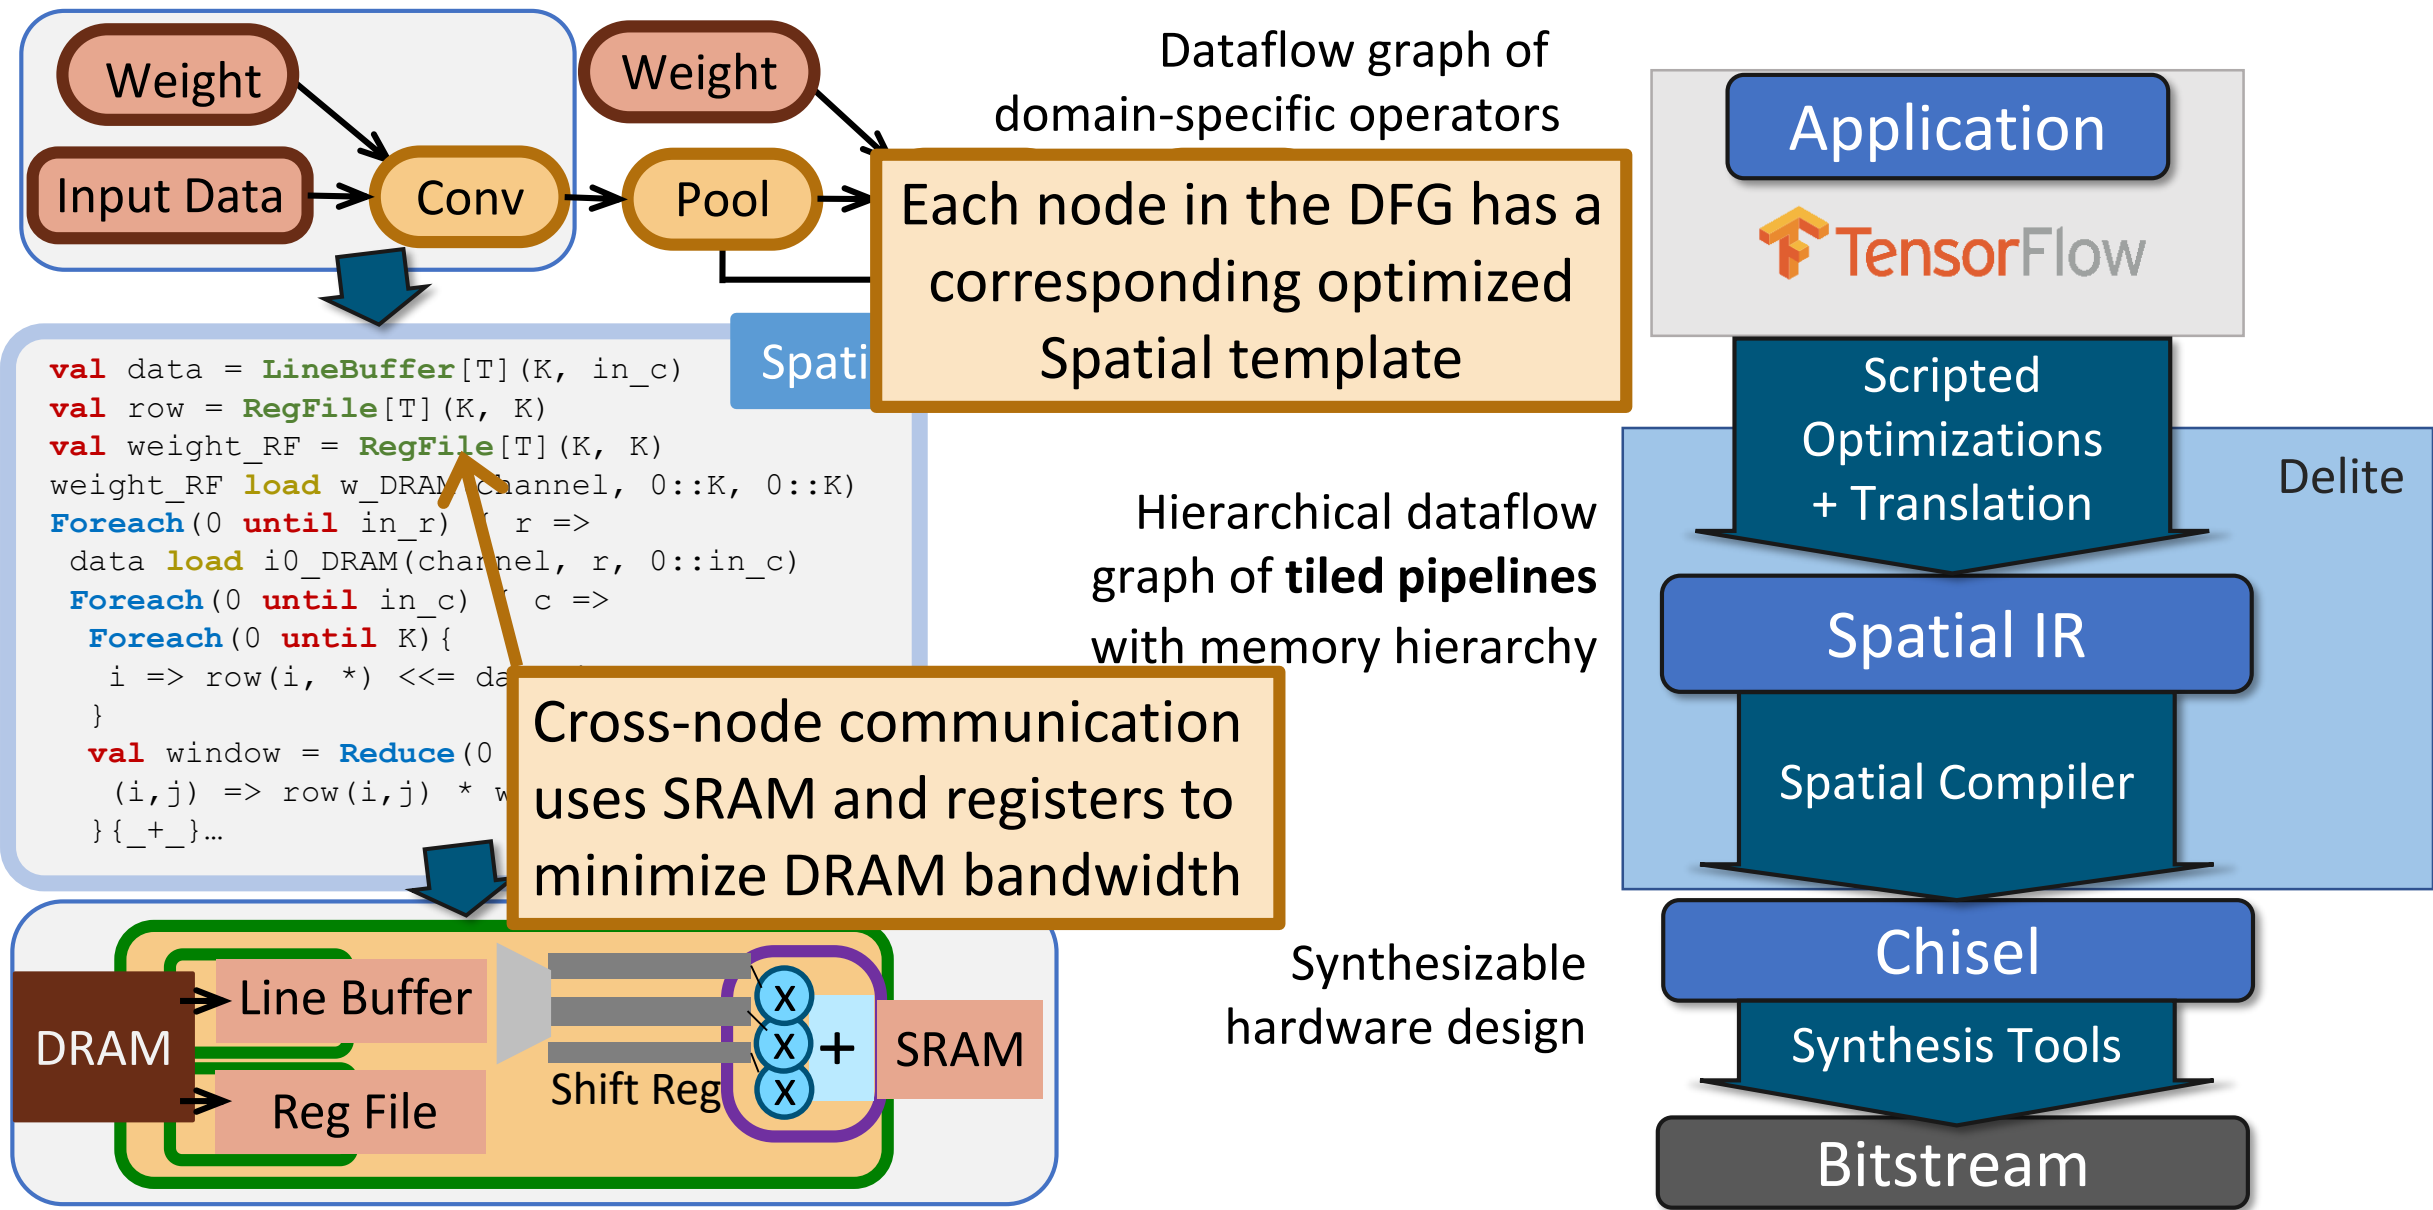
Space for GDA spans
four orders of magnitude

Resource Usage (% of maximum)

- Valid design point
- Invalid design point

- Pareto-optimal (ALMs/cycles) design
- Synthesized pareto design point

Preview: TensorFlow to Spatial



Preview: TensorFlow to Spatial

- Benefits

- Makes Spatial easier to program
- Automatically benefit from TensorFlow optimizations (e.g. constant folding, fusion, quantization for low precision)

- Current support:

- Convolutional Neural Networks
- Recurrent Neural Networks (GRU, LSTM)
- General matrix algebra computations

- Same DFG mapping can be applied to other front-ends

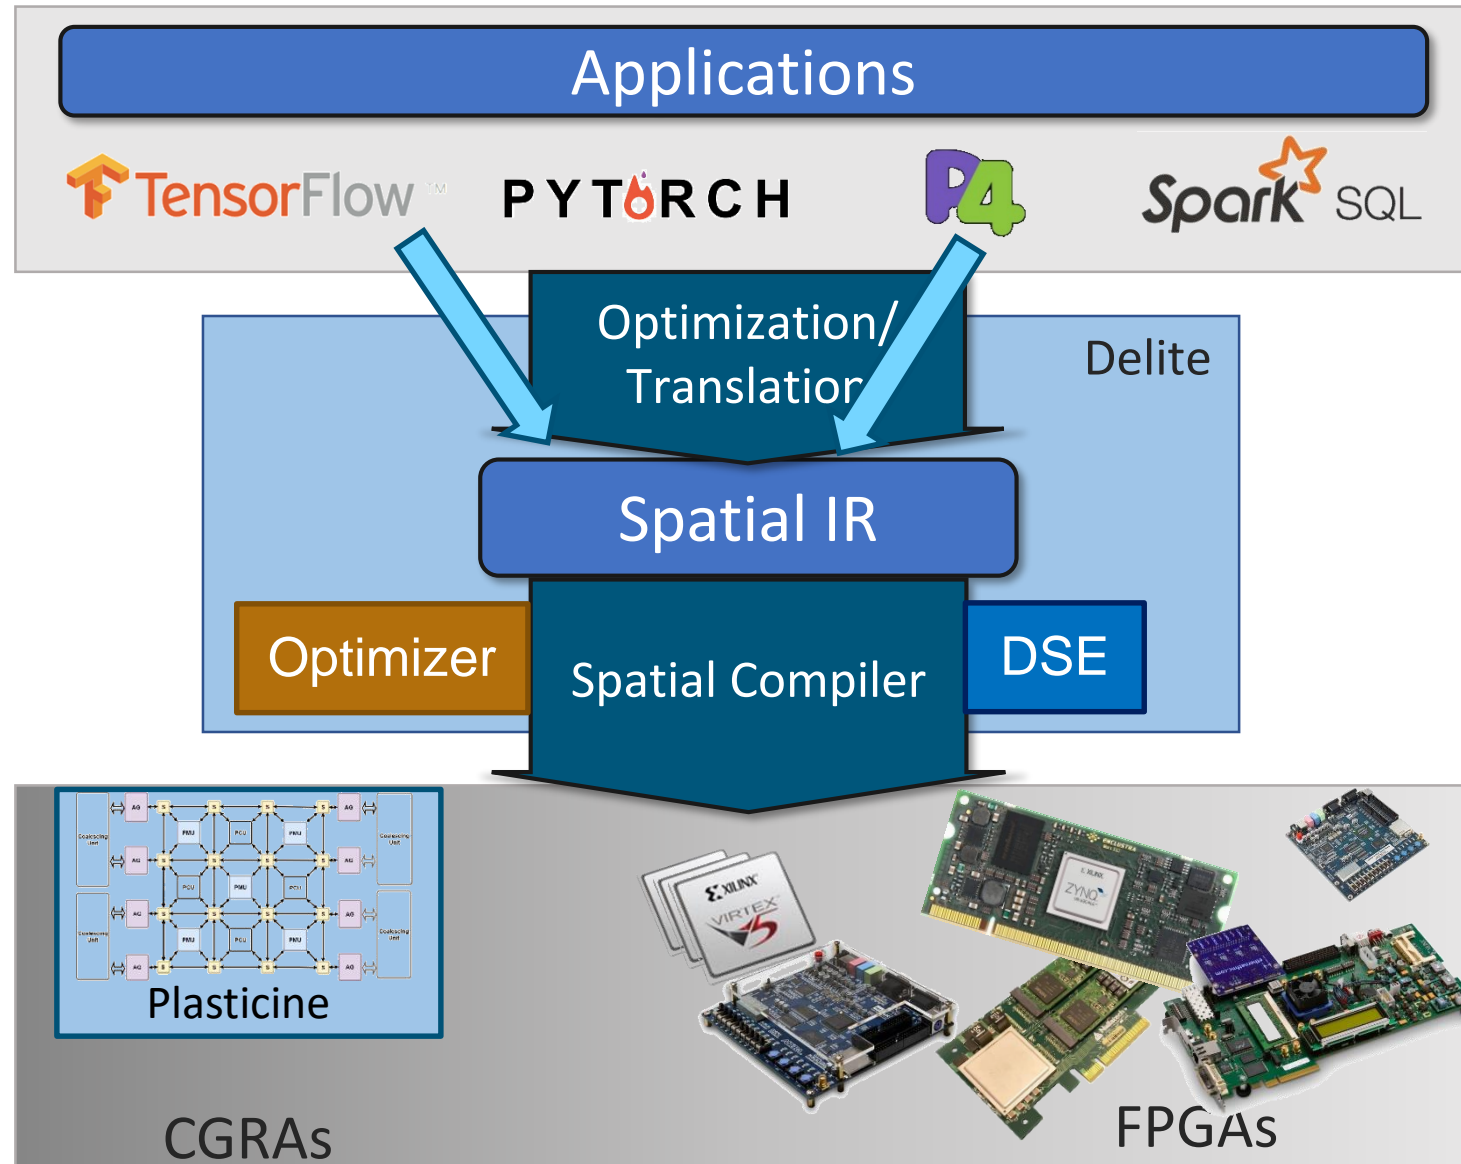
- Not specific to TensorFlow

Spatial's Advantages

- Automatic Design-Space Exploration
 - Parallelization selection to balance pipelines
 - Tile size selection to balance on-chip memory, bandwidth
- Automatic memory banking
- Agnostic to hardware vendor, Cloud/Embedded, FPGA/CGRA
- Preliminary results: within 20% of DNNWeaver (DNN hardware library)
 - For inferences/s of a small network (LeNet) on the Zynq ZC702
 - Ongoing work to close 20% gap by further optimizing templates

Future Work

- Low precision DNN training
- Residual Networks
- Unify and support more frontends above Spatial
- Expand set of optimizations done by Spatial compiler
- Expand backends supported by Spatial





Conclusion

- **Orders of magnitude improvements** in performance and energy efficiency are possible with reconfigurable architectures
- Spatial's **intermediate representation** can be used as a bridge to reconfigurable architectures
- The Spatial **language** includes abstractions for **performance, productivity, and portability**
- Spatial's design tuning is **~6500x** faster than that possible with Vivado HLS
- Preliminary results **within 20% of manual HDL** on CNN implementations

We're always looking for new users and feedback!

Open source at: spatial.stanford.edu



Nuts and Bolts of a Full Spatial Application

- Continue to see the bird's eye view of how to write a full application
- Visit the website tutorial for more detailed examples:
 - <http://spatial-lang.readthedocs.io/en/latest/tutorial/starting.html>

Spatial App Template

```
1 import spatial._
2 import org.virtualized._
3
4 object AppName extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8
9     ARM Host Code (setup)
10
11     Accel {
12       FPGA Code
13     }
14
15     ARM Host Code (teardown)
16   }
17 }
18
```

Spatial App Template

```
1 import spatial._
2 import org.virtualized._
3
4 object AppName extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8
9     - Define FPGA peripherals
10    - Send data from ARM to FPGA
11
12    Accel {
13      - Define FPGA operations
14    }
15    - Get data from the FPGA
16  }
17 }
18
```

Hello Spatial!

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8     val input = args(0).to[Int]
9     val in = ArgIn[Int]
10    val out = ArgOut[Int]
11    setArg(in, input)
12    Accel {
13      out := in + 4
14    }
15    println("Output: " + getArg(out))
16  }
17 }
18
```

Spatial is Embedded in Scala

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8     val input = args(0).to[Int]
9     val in = ArgIn[Int]
10    val out = ArgOut[Int]
11    setArg(in, input)
12    Accel {
13      out := in + 4
14    }
15    println("Output: " + getArg(out))
16  }
17 }
18
```

Spatial can be thought of as a **Scala** library

Spatial is Embedded in Scala

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8     val input = args(0).to[Int]
9     val in = ArgIn[Int]
10    val out = ArgOut[Int]
11    setArg(in, input)
12    Accel {
13      out := in + 4
14    }
15    println("Output: " + getArg(out))
16  }
17 }
18
```

Semicolons are optional

Import Statements

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6
7
8
9
10   val out := ArgOut[Int]
11   setArg(in, input)
12   Accel {
13     out := in + 4
14   }
15   println("Output: " + getArg(out))
16 }
17 }
18
```

Same in every Spatial program
(Similar idea to **#include** in C,
Identical to **import** in Java, Python)

Import Statements

```
1 import spatial._
```

Spatial-specific classes
(primarily **SpatialApp**)

```
2 import org.virtualized._
```

Useful macros for nicer
syntax (more later)

```
3 object HelloSpatial extends SpatialApp {  
4  
5   @virtualize  
6   def main(): Unit = {  
7     val input = args(0).to[Int]  
8
```

Application Object Declaration

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6     @Virtualize
7     def main(): Unit = {
```

Spatial applications are always **objects**

```
10     val out = ArgOut[Int]
11     setArg(in, input)
12     Accel {
13         out := in + 4
14     }
15     println("Output: " + getArg(out))
16 }
17 }
18
```

Application Object Declaration

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6     @virtualize
7     def main(): Unit = {
8
9         // ...
10        val out = ArgOut[Int]
11        setArg(in, input)
12        Accel {
13            out := in + 4
14        }
15        println("Output: " + getArg(out))
16    }
17 }
18
```

Name of application

All Spatial applications inherit from ("extends") **SpatialApp**

“@virtualize” Annotation

All functions in Spatial should have this annotation
(Allows overloading Scala constructs like if-then-else)

```
1 import spatial.  
5  
6 @virtualize  
7 def main(): Unit = {  
8     val input = args(0).to[Int]  
9     val in  = ArgIn[Int]  
10    val out = ArgOut[Int]  
11    setArg(in, input)  
12    Accel {  
13        out := in + 4  
14    }  
15    println("Output: " + getArg(out))  
16 }  
17 }  
18 }
```

Spatial's Entry Function: "main()"

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8     val input = args(0).toInt
9     val in = ArgIn[Int]
10    val out = ArgOut[Int]
11    setArg(in, input)
12    Accel {
13      out := in + 4
14    }
15    println("Output: " + getArg(out))
16  }
17 }
18
```

Spatial's entry function

Spatial's Entry Function: "main()"

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8     val input = args(0).to[Int]
9     val in = ArgIn[Int]
10    = ArgOut[Int]
11    n, input
12    out := in + 4
13  }
14
15   println("Output: " + getArg(out))
16 }
17 }
18 }
```

Starts a function declaration

Function return type (Unit: same as void)

Val Definitions

Declares an **immutable** value named “input” (value can’t be modified later)

```
1 import spatial
2
3
4
5
6 @virtualize
7 def main(): Unit = {
8     val input = args(0).to[Int]
9     val in    = ArgIn[Int]
10    val out = ArgOut[Int]
11    setArg(in, input)
12    Accel {
13        out := in + 4
14    }
15    println("Output: " + getArg(out))
16 }
17 }
18 }
```


Val Definitions


Value types are optional in Scala.

```
1 import spatial._
2 import org.virtualized._
3
4
5
6 @virtualize
7 def main(): Unit = {
8     val input: Int = args(0).to[Int]
9     val in  = ArgIn[Int]
10    val out = ArgOut[Int]
11    setArg(in, input)
12    Accel {
13        out := in + 4
14    }
15    println("Output: " + getArg(out))
16 }
17 }
18 }
```

Val Definitions

Scala is statically typed (like C, Java)
Without the “: **Int**”, the type of this value is **inferred** by the compiler.

```
5  
6 @virtualize  
7 def main(): Unit = {  
8     val input = args(0).to[Int]  
9     val in = ArgIn[Int]  
10    val out = ArgOut[Int]  
11    setArg(in, input)  
12    Accel {  
13        out := in + 4  
14    }  
15    println("Output: " + getArg(out))  
16 }  
17 }  
18 }
```

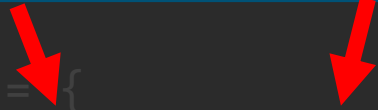


Method Calls

```
1 import spatial._
```

Round brackets () for value parameters
Square brackets [] are for **type** parameters


```
6  @virtualize
7  def main(): Unit = {
8      val input = args(0).to[Int]
9      val in  = ArgIn[Int]
10     val out = ArgOut[Int]
11     setArg(in, input)
12     Accel {
13         out := in + 4
14     }
15     println("Output: " + getArg(out))
16 }
17 }
18 }
```



Spatial Command-Line Arguments

Spatial app's command-line arguments

Conversion from **String** to **Int**




```
int main(int argc, char **argv) {  
    int in = atoi(argv[1]);  
  
    printf("Output: %d\n", out);  
    return 0;  
}
```

Input Arguments (ArgIn)

Creates a new register to capture a scalar argument *from* the ARM


```
1 import spatial._
2 import org.virtualized._
3
4
5
6
7 def main(): Unit = {
8   val input = args(0).to[Int]
9   val in = ArgIn[Int]
10  val out = ArgOut[Int]
11  setArg(in, input)
12  Accel {
13    out := in + 4
14  }
15  println("Output: " + getArg(out))
16 }
17 }
18 }
```



Output Arguments (ArgOut)

Creates a new scalar argument *to* the ARM *from* the FPGA


```
1 import spatial._
2 import org.virtualized._
3
4 @HiFiHelloSpatial extends SpatialApp {
5
6   val input = args(0).to[Int]
7   val in = ArgIn[Int]
8   val out = ArgOut[Int]
9   setArg(in, input)
10   Accel {
11     out := in + 4
12   }
13   println("Output: " + getArg(out))
14 }
15
16 }
17
18 }
```



Scalar Transfers (ARM → FPGA)

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8     val input = args(0).to[Int]
9     val in = ArgIn[Int]
10    val out = ArgOut[Int]
11    setArg(in, input)
12    Accel {
13
14
15
16
17 }
18
```

Tells the host ARM to write **input** to scalar argument **in** on the FPGA




Accel Block

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
```

Defines an FPGA computation scope.
Everything in here goes on the FPGA


```
10   val out = ArgOut[Int]
11   setArg(in, input)
12   Accel {
13     out := in + 4
14   }
15   println("Output: " + getArg(out))
16 }
17 }
18 }
```



Accel Block

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6
7
8
9
10
11   setArg(in, input)
12   Accel {
13     out := in + 4
14   }
15   println("Output: " + getArg(out))
16 }
17 }
18 }
```

The types of operations that can be done in this scope are limited to **synthesizable** Spatial




Accel Block

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
```

Accel handles control signals for you.
It implicitly creates:


- a **start signal** (ARM \rightarrow FPGA)
- a **done signal** (FPGA \rightarrow ARM)

```
11   setArg(in, input)
12   Accel {
13     out := in + 4
14   }
15   println("Output: " + getArg(out))
16 }
17 }
18
```




Implicit Register Reads

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main() {
8
9     // Implicitly creates a wire from the
10    // register (ArgIn) in
11    setArg(in, input)
12    Accel {
13      out := in + 4
14    }
15    println("Output: " + getArg(out))
16  }
17 }
18
```




Register Writes

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8
9     := creates a write of the value
10    in + 4 to the register out
11    setArg(in, input)
12    Accel {
13      out := in + 4
14    }
15    println("Output: " + getArg(out))
16  }
17 }
18
```



Accel Block Scheduling

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   Accel guarantees that FPGA
7   execution completes after all
8   operations in this block complete
9
10   val out = ArgOut[Int]
11   setArg(in, input)
12   Accel {
13     out := in + 4
14   }
15   println("Output: " + getArg(out))
16 }
17 }
18 }
```




Scalar Transfers (FPGA → ARM)

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8     val input = args(0).to[Int]
9
10    

Gets the value of the ArgOut out from the FPGA back to the ARM



11
12    out := in + 4
13  }
14  println("Output: " + getArg(out))
15 }
16
17 }
18
```



Printing in Spatial**

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8     val input = args(0).to[Int]
9     val in = ArgIn[Int]
10    val out = ArgOut[Int]
11
12    // Prints the output to the terminal
13    out := in + 4
14  }
15  println("Output: " + getArg(out))
16 }
17
18
```

** Printing in Spatial isn't synthesizable, but it can be used in **host code** and in **debugging** (more in future lectures)



```
int main(int argc, char **argv) {
  int in = atoi(argv[1]);
  ...
  printf("Output: %d\n", out);
  return 0;
}
```

Hello Spatial!

```
1 import spatial._
2 import org.virtualized._
3
4 object HelloSpatial extends SpatialApp {
5
6   @virtualize
7   def main(): Unit = {
8     val input = args(0).to[Int]
9     val in = ArgIn[Int]
10    val out = ArgOut[Int]
11    setArg(in, input)
12    Accel {
13      out := in + 4
14    }
15    println("Output: " + getArg(out))
16  }
17 }
18
```


Custom Types in Spatial

- Now what if we want an ArgIn value that isn't an Int?
- Other options:
 - Custom fixed point types
 - Custom floating point types
 - Structs
 - Vectors

Custom Types

```
1
2
3 val input = args(0).to[Int]
4 val in = ArgIn[Int]
5
6 setArg(in, input)
```

Custom Fixed Point Types

```
1 type Q8_8 = FixPt[FALSE, _8, _8]
```

_N = # of fraction bits
(N from 0 to 128)

TRUE = Signed
FALSE = Unsigned

_N = # of integer bits
(N from 1 to 128)

```
11 0b00000000.00000000
```

12 └───┘ └───┘

13 Integer bits Fraction bits

Custom Fixed Point Examples

```
1 type Q8_8 = FixPt[FALSE,_8,_8]
2
3 type UInt8 = FixPt[FALSE,_8,_0]
4
5 type LongLong = FixPt[TRUE,_128,_0]
```

0b00000000.00000000

Integer bits Fraction bits

Custom Fixed Point Types

```
1 type UInt8 = FixPt[FALSE,_8,_0]
2
3 val input = args(0).to[UInt8]
4 val in  = ArgIn[UInt8]
5
6 setArg(in, input)
```

Custom Floating Point Types

```
1 type Float = FltPt[_23,_11]
```

_N = # of significand bits + 1
(N from 1 to 128)
Includes sign bit!

_N = # of exponent bits
(N from 0 to 128)

```
10 0 00000000 x 2^00000000
```

Sign bit Significand bits Exponent bits

Custom Floating Point Types

```
1 type Half = FltPt[_11,_5]
2
3 val input = args(0).to[Half]
4 val in  = ArgIn[Half]
5
6 setArg(in, input)
```

Predefined Type Aliases

```
1 type Char  = FixPt[TRUE,_8,_0]
2 type Short = FixPt[TRUE,_16,_0]
3 type Int   = FixPt[TRUE,_32,_0]
4 type Long  = FixPt[TRUE,_64,_0]
5
6 type Half   = FltPt[_11,_5]    // 754 Half
7 type Float  = FltPt[_24,_8]    // 754 Single
8 type Double = FltPt[_53,_11]   // 754 Double
9
10
11
12
13
14
15
16
17
18
```


Note About Booleans

```
1
2
3 val input = args(0).to[Boolean]
4 val in   = ArgIn[Boolean]
5
6 setArg(in, input)
```

Note: For API purposes,
Boolean is NOT the same as
single bit fixed point number


Uses “false” and “true”
rather than 0 and 1

Custom Structs

```
1 @struct class MyStruct(  
2     red:    Int,  
3     green:  Int,  
4     blue:   Int  
5 )  
6
```

Declares a new Struct type
with the given list of fields

Custom Structs




```
1 @struct class MyStruct(  
2   red:  Int,  
3   green: Int,  
4   blue:  Int  
5 )  
6  
7 val bus = MyStruct(147, 192, 125)
```

Allocates an instance of the struct.
Note: NO *new* keyword used

In hardware, a struct instance is just
a concatenation of wires


Custom Structs



```
1 @struct class MyStruct(  
2     red:  Int,  
3     green: Int,  
4     blue:  Int  
5 )  
6  
7 val bus = MyStruct(147, 192, 125)  
8  
9 val red = bus.red  
10 val blue = bus.blue
```

Creates a reference to the struct field (equivalent to a bit slice)

Custom Structs



```
1 @struct class MyStruct(  
2     red:  Int,  
3     green: Int,  
4     blue:  Int  
5 )  
6  
7 val bus = MyStruct(147, 192, 125)
```

```
8  
9 bus.blue = 45
```

Note: Allocated structs are immutable!
We can't write to them or change the contents!

Nesting Structs

```
1 @struct class RGB(  
2     red:  Int,  
3     green: Int,  
4     blue: Int  
5 )  
6  
7 @struct class RGBA(  
8     rgb:  RGB,  
9     alpha: Int  
10 )  
11  
12  
13  
14  
15  
16  
17  
18
```

Registers of Custom Types

```
1 @struct class MyStruct(  
2     red:    Int,  
3     green:  Int,  
4     blue:   Int  
5 )  
6  
7 val in  = ArgIn[MyStruct]  
8  
9 in.red
```

Creates an ArgIn register which holds a value of type MyStruct

Note: Registers can hold structs as long as the fields are primitive values (FixPt, FltPt, Boolean) or other primitive-based structs